

**CS3500 Operating Systems**  
**Course Project: Address Space Layout**  
**Randomisation**

**Submitted by:**  
CS18B025 M Harini Saraswathy  
CS18B045 Rishika Varma K

December 5, 2020

## **Introduction:**

Address Space Layout Randomisation (ASLR) is a security technique in which the address at which executable, stack, heap, libraries, etc are loaded at randomised locations as opposed to being loaded at the same fixed place. This technique helps in preventing buffer overflow attacks as the address at which each part of code and data is unknown and changes each time it is run.

ASLR increases the control-flow integrity of a system by making it more difficult for an attacker to execute a successful buffer-overflow attack by randomizing the offsets it uses in memory layouts. ASLR is used today on Linux, Windows, and MacOS systems. In this project, we implemented ASLR in the xv6 operating system.

## **Advantages and Disadvantages of ASLR:**

### **Advantages:**

ASLR (Address Space Layout Randomization) is a memory exploitation mitigation technique used on both Linux and Windows systems. Address Space Layout Randomization (ASLR) is a memory-protection process for operating systems that guards against buffer-overflow attacks. It helps to ensure that the memory addresses associated with running processes on systems are not predictable, thus flaws or vulnerabilities associated with these processes will be more difficult to exploit.

### **Disadvantages:**

The effectiveness of ASLR is dependent on the entirety of the address space layout remaining unknown to the attacker. In addition, only executables that are compiled as Position Independent Executable (PIE) programs will be able to claim the maximum protection from ASLR technique because all sections of the code will be loaded at random locations.

## **Modifications required to incorporate ASLR into xv6**

### **Randomisation function**

#### **Approach:**

The first step is implementing the randomise function that is required for generating random numbers which are used as the offsets to the addresses at which stack, executable, etc are to be loaded. This was implemented using the Lehmer random number generator algorithm (Source: Wikipedia). The algorithm which was implemented is shown below.

```

uint32_t lcg_parkmiller(uint32_t *state)
{
    // Precomputed parameters for Schrage's method
    const uint32_t M = 0x7fffffff;
    const uint32_t A = 48271;
    const uint32_t Q = M / A;    // 44488
    const uint32_t R = M % A;    // 3399

    uint32_t div = *state / Q;    // max: M / Q = A = 48,271
    uint32_t rem = *state % Q;    // max: Q - 1 = 44,487

    int32_t s = rem * A;    // max: 44,487 * 48,271 = 2,147,431,977 = 0x7fff3629
    int32_t t = div * R;    // max: 48,271 * 3,399 = 164,073,129
    int32_t result = s - t;

    if (result < 0)
        result += M;

    return *state = result;
}

```

## Implementation:

We implemented the above function in the file sysfile.c and the initial value of the seed used was the value of ticks. We have been kept track of this value in a variable hashticks using the ticks value in trap.c.

## Code:

```

485 static unsigned random_seed = 1;
486 static int present = 0;
487
488 unsigned lcg_parkmiller()
489 {
490     const unsigned N = 0x7fffffff;
491     const unsigned G = 48271u;
492     unsigned div = random_seed / (N / G);
493     unsigned rem = random_seed % (N / G);
494
495     unsigned a = rem * G;
496     unsigned b = div * (N % G);
497
498     return random_seed = (a > b) ? (a - b) : (a + (N - b));
499 }
500
501 uint64 random(int start, int end) {
502     if(present == 1){
503         return (lcg_parkmiller() % (end-start))+start;
504     }
505     random_seed = hashticks;
506     present = 1;
507     return (lcg_parkmiller() % (end-start))+start;
508 }

```

## Adding support to enable aslr for a particular process

### Approach:

This was achieved by adding a syscall that enables a variable signifying usage of aslr and an aslr.c file which can be called from user space. For a process to be run with aslr then we must call the aslr process with 2 arguments, the first being the name of the process to be run with aslr and the second argument is the argv for that process.

### Implementation:

We added an aslr\_var field in the proc structure to signify whether process is run with aslr. We added a syscall called aslr which sets the aslr\_var field in proc which was initially 0 to 1. The aslr\_var field in the proc structure is initialised to 0 in allocproc() (the constructor) and freeproc() (the destructor). We created a file aslr.c that calls the aslr syscall which in turn enables the aslr\_var field in the proc for that process. After it returns from the syscall, exec is called on the process to be run whose name is given as the first argument to the aslr process and the remaining part of argv apart from the first string as the argv for the process (second argument in exec) . Thus aslr can be enabled for a particular process by calling aslr process with arguments as the name of the process to be run and its arguments.

### Code:

```
99     uint64
100     sys_aslr(void)
101     {
102         myproc()->aslr_var=1;
103         return 0;
104     }
105
```

```
1  #include "kernel/types.h"
2  #include "kernel/stat.h"
3  #include "user/user.h"
4
5  int
6  main(int argc, char *argv[])
7  {
8      aslr();
9      exec(argv[1], (argv+1));
10     exit(0);
11 }
```

## Incorporating loading of executable at random offset

### Approach:

We generate a random offset using the randomize function and start loading the executable from that address. As it is randomly generated it may not be page aligned and so the implementation of loadseg must be appropriately changed to take care of this issue.

### Implementation:

A variable load\_offset is maintained which is 0 if en\_aslr is disabled and a random number between 0 to 1000 left shifted by 4 if it is enabled. This is the offset at which the executable and relocation addresses are shifted by. The loadseg function is then called with this offset added to the virtual address. In the loadseg function, the starting part of the segment which is not page aligned is loaded separately according to the size in a way similar to how the ending part of the segment that is not page aligned is loaded, and the rest of the segment is loaded through the for loop as was the case earlier (no modification in that aspect).

### Code:

```
60     if(en_aslr)
61         load_offset = random(0, 1000) << 4;
62     else load_offset = 0;
63     //load_offset = PGROUNDUP(load_offset);
64     printf("load offset: 0x%x\n", load_offset);
65
66     // Read in ELF PROG_LOAD programs into memory
67     sz = uvmmalloc(pagetable, 0, load_offset);
68     for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){
69         if(readi(ip, 0, (uint64)&ph, off, sizeof(ph)) != sizeof(ph)) {
70             printf("exec: readi error\n");
71             goto bad;
72         }
73         if(ph.type != ELF_PROG_LOAD)
74             continue;
75         if(ph.memsz < ph.filesz) {
76             printf("exec: memsz smaller than filesz error\n");
77             goto bad;
78         }
79         if((sz = uvmmalloc(pagetable, sz, ph.vaddr + ph.memsz + load_offset)) == 0) {
80             printf("exec: uvmmalloc error\n");
81             goto bad;
82         }
83         if(loadseg(pagetable, ph.vaddr + load_offset, ip, ph.off, ph.filesz) < 0) {
84             printf("exec: loadseg error\n");
85             goto bad;
86         }
87     }
```

```

193 // Load a program segment into pagetable at virtual address va.
194 // and the pages from va to va+sz must already be mapped.
195 // Returns 0 on success, -1 on failure.
196 static int
197 loadseg(pagetable_t pagetable, uint64 va, struct inode *ip, uint offset, uint sz)
198 {
199     uint i, n;
200     uint64 pa;
201
202     // get the positive page offset of the va from the beginning of the page
203     uint64 first_va = PGROUNDDOWN(va);
204     uint64 pg_offset = va - first_va;
205
206     // manually fill the first page, which might not be page aligned
207     pa = walkaddr(pagetable, first_va);
208     if (pa == 0)
209         panic("loadseg: address should exist");
210     // fill page with zeroes
211     memset((void*)pa, 0, PGSIZE);
212     // fill rest of page
213     n = (sz < PGSIZE - pg_offset)? sz : PGSIZE - pg_offset;
214     if(readi(ip, 0, (uint64)pa + pg_offset, offset, n) != n)
215         return -1;
216     offset += n;
217     sz -= n;
218
219     // use for loop for remaining pages
220     for(i = PGSIZE; sz > 0; i += PGSIZE){
221         pa = walkaddr(pagetable, first_va + i);
222         if(pa == 0)
223             panic("loadseg: address should exist");
224         // zero the page
225         memset((void*)pa, 0, PGSIZE);
226         // fill the page or until there are no bytes left to write
227         n = (sz < PGSIZE)? sz : PGSIZE;
228         if(readi(ip, 0, (uint64)pa, offset, n) != n)
229             return -1;
230         offset += n;
231         sz -= n;
232     }
233
234     return 0;
235 }
236

```

## Incorporation of changing corresponding relocation offsets

### Approach:

The relocation section header consists of the offsets and other information at which relocatable attributes are loaded. As the executable has been loaded at an offset the corresponding relocation offsets also need to be modified. This is done by going through all the section headers till we reach the relocation section header and then we increase the offset of the corresponding entries to be changed by load\_offset value.

### Implementation:

We created 2 extra structures in elf.h corresponding to section header and relocation entries respectively. These structures were used to read each section header from the elfhdr struct to check which of them is the relocations section header and to read each relocation entry from the section header from which the type of entry is read and the offset is changed (increased by load\_offset) correspondingly if required.

## Code:

```
89 // Get Section Headers
90 for(i=0, off = elf.shoff; i < elf.shnum; i++, off += elf.shentsize) {
91     if (readi(ip, 0, (uint64)&sh, off, elf.shentsize) != elf.shentsize)
92         goto bad;
93
94     int nr = (sh.type ^ 4);
95     if (!nr) {
96         // found section header for relocations
97         // read through each relocation
98         for (int sectoff = 0, relocnum = 1; sectoff < sh.size; sectoff += sh.entsize, relocnum++) {
99             int size = readi(ip, 0, (uint64)&reloc, sh.offset + sectoff, sh.entsize);
100
101             if((reloc.info&0xffffffffL) == 2 || (reloc.info&0xffffffffL) == 3 || (reloc.info&0xffffffffL) == 5){
102                 if (copyin(pagetable, (char*)&instr, (uint64)reloc.offset + load_offset, 8) != 0)
103                     goto bad;
104                 instr += load_offset;
105                 if (copyout(pagetable, (uint64)reloc.offset + load_offset, (char*)&instr, 8) != 0)
106                     goto bad;
107             }
108
109             if (size != sizeof(struct elfrel))
110                 goto bad;
111         }
112     }
113 }
```

## Adding support for stack offset

### Approach:

The other aspect of aslr is initialising stack at a random address as well. A random number of pages are left empty at the start after which the stack page is made.

### Implementation:

A variable `stack_offset` is created and initialised to 2 if aslr is disabled or to a random number between 2 and 1000 obtained using `randomize` function if aslr is enabled. A total memory of the required size of the process that has been kept track of (rounded up to page align) and `stack_offset` number of pages is allocated. `sp` is then initialised to the size of memory that has been initialised (which is effectively the address just after the end of the stack page) and `stackbase` as the address of the start of the stack page (`sp - page size`). The rest of the exec process remains same as earlier in regards to creating new page directory, trapframe, etc.

## Code:

```
121 // Allocate random number of pages at the next page boundary.
122 // Use the last one as the user stack.
123 if(en aslr)
124     stack_offset = random(2, 1000);
125 else
126     stack_offset = 2;
127 sz = PGROUNDUP(sz);
128 if((sz = uvmmalloc(pagetable, sz, sz + stack_offset*PGSIZE)) == 0)
129     goto bad;
130 uvmmclear(pagetable, sz-stack_offset*PGSIZE);
131 sp = sz;
132 stackbase = sp - PGSIZE;
133
134 // Push argument strings, prepare rest of stack in ustack.
135 for(argc = 0; argv[argc]; argc++) {
136     if(argc >= MAXARG)
137         goto bad;
138     sp -= strlen(argv[argc]) + 1;
139     sp -= sp % 16; // riscv sp must be 16-byte aligned
140     if(sp < stackbase)
141         goto bad;
142     if(copyout(pagetable, sp, argv[argc], strlen(argv[argc]) + 1) < 0)
143         goto bad;
144     ustack[argc] = sp;
145 }
146 ustack[argc] = 0;
```

## Modifications made to show proof of working

To show that the technique is working we printed certain values in `exec.c`. The values that will change are the place at which executable is loaded, in which case the entry point at which the implementation enters (main function) changes for each time the process is run. Thus effectively the initial entry point address is stored in the `epc` register and so this is printed to check if it varies. Similarly, the other thing that has been randomized is the starting address of the stack which is the same as the initial value of `sp` and so to show the randomisation in `sp` this is also printed. An extra message is printed to indicate whether `aslr` is enabled or disabled.

### Code:

```
161 // Save program name for debugging.
162 for(last=s=path; *s; s++)
163     if(*s == '/')
164         last = s+1;
165 safestrcpy(p->name, last, sizeof(p->name));
166
167 // Commit to the user image.
168 oldpagetable = p->pagetable;
169 p->pagetable = pagetable;
170 p->sz = sz;
171 p->tf->epc = elf.entry + load_offset; // initial program counter = main
172 p->tf->sp = sp; // initial stack pointer
173 proc_freepagetable(oldpagetable, oldsz);
174
175 printf("sp: 0x%x\n", r_sp());
176 printf("epc: 0x%x\n", p->tf->epc);
177 if(en_aslr)
178     printf("ASLR is enabled hence epc changes with each run of aslr.\n");
179 else
180     printf("ASLR is disabled hence epc remains constant.\n");
181 return argc; // this ends up in a0, the first argument to main(argc, argv)
182
183 bad:
184 if(pagetable)
185     proc_freepagetable(pagetable, sz);
186 if(ip){
187     iunlockput(ip);
188     end_op(ROOTDEV);
189 }
190 return -1;
191 }
```

### Code of the process that is tested:

```
1 #include "kernel/types.h"
2 #include "user/user.h"
3 #include "kernel/fcntl.h"
4
5 int temp() {
6     printf("Tester function for ASLR\n");
7     return 1;
8 }
9
10 int main()
11 {
12     if(temp() != 1)
13         exit(1);
14
15     int fds[2];
16
17     if(pipe(fds) != 0){
18         printf("pipe() failed\n");
19         exit(1);
20     }
21     else printf("pipe() passed\n");
22
23     int pid = fork();
24     printf("fork() passed : %d\n", pid);
25
26     printf("Congrats! Qemu lives! \n");
27     exit(0);
28 }
29
```




Output:

Aslr disabled:

The screenshot shows a terminal window with the following content:

```
Activities Terminal
File Edit View Search Terminal Help
virtio disk init 0
exec /init:
ASLR Disabled!
load offset: 0x0
sp: 0x-2480
epc: 0x0
ASLR is disabled hence epc remains constant.
init: starting sh
exec sh:
ASLR Disabled!
load offset: 0x0
sp: 0x-4480
epc: 0xa00
ASLR is disabled hence epc remains constant.
$ t
exec t:
ASLR Disabled!
load offset: 0x0
sp: 0x-6480
epc: 0x22
ASLR is disabled hence epc remains constant.
Tester function for ASLR
pipe() passed
fork() passed : 4
Congrats! Qemu lives!
fork() passed : 0
Congrats! Qemu lives!
$ t
exec t:
ASLR Disabled!
load offset: 0x0
sp: 0x-6480
epc: 0x22
ASLR is disabled hence epc remains constant.
Tester function for ASLR
pipe() passed
fork() passed : 6
Congrats! Qemu lives!
fork() passed : 0
Congrats! Qemu lives!
$ t
exec t:
ASLR Disabled!
load offset: 0x0
sp: 0x-6480
epc: 0x22
ASLR is disabled hence epc remains constant.
Tester function for ASLR
pipe() passed
fork() passed : 8
Congrats! Qemu lives!
fork() passed : 0
Congrats! Qemu lives!
$ aslr t
exec aslr:
```

Aslr enabled:



```

Activities Terminal Thu 01:15 harini@harini-saraswathy: ~/ASLR
File Edit View Search Terminal Help
ASLR Disabled!
load offset: 0x0
sp: 0x-6480
pc: 0x0
ASLR is disabled hence epc remains constant.
exec t;
ASLR Enabled!
load offset: 0x2cd0
sp: 0x-6480
pc: 0x2cf2
ASLR is enabled hence epc changes with each run of aslr.
Tester function for ASLR
pipe() passed
fork() passed : 10
Congrats! Qemu lives!
fork() passed : 0
Congrats! Qemu lives!
$ aslr t;
exec aslr:
ASLR Disabled!
load offset: 0x0
sp: 0x-6480
pc: 0x0
ASLR is disabled hence epc remains constant.
exec t;
ASLR Enabled!
load offset: 0x3cf0
sp: 0x-6480
pc: 0x3d12
ASLR is enabled hence epc changes with each run of aslr.
Tester function for ASLR
pipe() passed
fork() passed : 12
Congrats! Qemu lives!
fork() passed : 0
Congrats! Qemu lives!
$ aslr t;
exec aslr:
ASLR Disabled!
load offset: 0x0
sp: 0x-6480
pc: 0x0
ASLR is disabled hence epc remains constant.
exec t;
ASLR Enabled!
load offset: 0x3360
sp: 0x-6480
pc: 0x3382
ASLR is enabled hence epc changes with each run of aslr.
Tester function for ASLR
pipe() passed
fork() passed : 14
Congrats! Qemu lives!
fork() passed : 0
Congrats! Qemu lives!

```

## References:

- <https://nptel.ac.in/courses/106/106/106106199/>
- [https://en.wikipedia.org/wiki/Lehmer\\_random\\_number\\_generator](https://en.wikipedia.org/wiki/Lehmer_random_number_generator)
- [https://www.reddit.com/r/osdev/comments/8b4tnj/aslr\\_implementation/](https://www.reddit.com/r/osdev/comments/8b4tnj/aslr_implementation/)
- <https://iitd-plos.github.io/os/2020/ref/os-arpaci-dessau-book.pdf>
- <http://pages.cs.wisc.edu/~remzi/OSTEP/>
- <https://github.com/johnmwu/xv6-aslr>
- <https://github.com/tenkjm/xv6-aslr>
- <https://github.com/typingkla/xv6-aslr>
- <https://www.networkworld.com/article/3331199/what-does-aslr-do-for-linux.html>
- <https://tc.gts3.org/cs3210/2016/spring/p/prop-slides.pdf>
- <https://blog.morphisec.com/aslr-what-it-is-and-what-it-isnt/>
- [https://docs.oracle.com/cd/E23824\\_01/html/819-0690/chapter6-54839.html](https://docs.oracle.com/cd/E23824_01/html/819-0690/chapter6-54839.html)
- <https://refspecs.linuxbase.org/elf/gabi4+/ch4.reloc.html>
- <https://stackoverflow.com/questions/16847741/processing-elf-relocations-understanding-the-relocs-symbols-section-data-an>
- [http://blog.k3170makan.com/2018/10/introduction-to-elf-format-part-vi\\_18.html](http://blog.k3170makan.com/2018/10/introduction-to-elf-format-part-vi_18.html)